

MESURER LES EMISSIONS DE CARBONE DES TRAITEMENTS ML



PROBLÉMATIQUE SOCIÉTALE

- ▶ Augmentation de la taille & complexité des modèles
- ▶ Augmentation de l'utilisation de modèles de ML
- ▶ Besoin d'un outil de mesure pour plusieurs buts :
 - ▶ Encourager la sobriété numérique
 - ▶ Point de vue Y. Bengio : Data Scientists doivent mesurer pour pouvoir mettre en place des taxes pigouviennes pour forcer les corporations à réduire leur empreinte carbone


Consumption	CO ₂ e (lbs)
Air travel, 1 passenger, NY↔SF	1984
Human life, avg, 1 year	11,023
American life, avg, 1 year	36,156
Car, avg incl. fuel, 1 lifetime	126,000
Training one model (GPU)	
NLP pipeline (parsing, SRL)	39
w/ tuning & experimentation	78,468
Transformer (big)	192
w/ neural architecture search	626,155

Table 1: Estimated CO₂ emissions from training common NLP models, compared to familiar consumption.¹



00:00:00:00


EMPREINTE CARBON DANS LES TRAITEMENTS ML



MESURER LA CONSOMMATION D'UNE MACHINE PHYSIQUE

- Métrique de mesure de la consommation d'un CPU : taux d'utilisation * Thermal Design Power

"the power consumption (in watts) under the maximum theoretical load. Power consumption is less than TDP under lower loads. The TDP is the maximum power that one should be designing the system for."



MESURER LA CONSOMMATION D'UNE MACHINE PHYSIQUE

- ▶ Mesure lue, dans les OS équipés de processeurs compatibles, dans le fichier **RAPL**
- ▶ Utilisation de l'Intel Power Gadget pour les processeurs sans interface **RAPL**
- ▶ Consommation de la mémoire également grâce au TDP & dans le fichier **RAPL**

```
cpu0: MSR_IA32_TEMPERATURE_TARGET: 0x005a0a00 (90 C)
cpu24: MSR_IA32_TEMPERATURE_TARGET: 0x005a0a00 (90 C)
cpu0: MSR_IA32_PACKAGE_THERM_STATUS: 0x88300000 (42 C)
cpu0: MSR_IA32_PACKAGE_THERM_INTERRUPT: 0x00000000 (90 C, 90 C)
cpu24: MSR_IA32_PACKAGE_THERM_STATUS: 0x88370000 (35 C)
cpu24: MSR_IA32_PACKAGE_THERM_INTERRUPT: 0x00000000 (90 C, 90 C)
cpu17: MSR_PKG_C3_IRTLL: 0x00000000 (NOTvalid, 0 ns)
cpu17: MSR_PKG_C6_IRTLL: 0x00000000 (NOTvalid, 0 ns)
cpu17: MSR_PKG_C7_IRTLL: 0x00000000 (NOTvalid, 0 ns)
Package Core CPU Avg_MHz Busy% Bzy_MHz TSC_MHz IRQ SMI CPU%c1 CPU%c6 CoreTmp PkgTmp Pkg%pc2 Pkg%pc6 PkgWatt RAMWatt PKG_% RAM_%
- - - 0 0.00 2164 3000 257 0 100.00 0.00 42 42 0.00 0.00 117.21 73.34 0.00 0.00
0 0 0 0.05 1783 3001 26 0 99.95 0.00 42 42 0.00 0.00 58.97 33.62 0.00 0.00
1 0 24 0 0.00 1200 3000 2 0 100.00 0.00 33 37 0.00 0.00 58.23 39.72 0.00 0.00
```

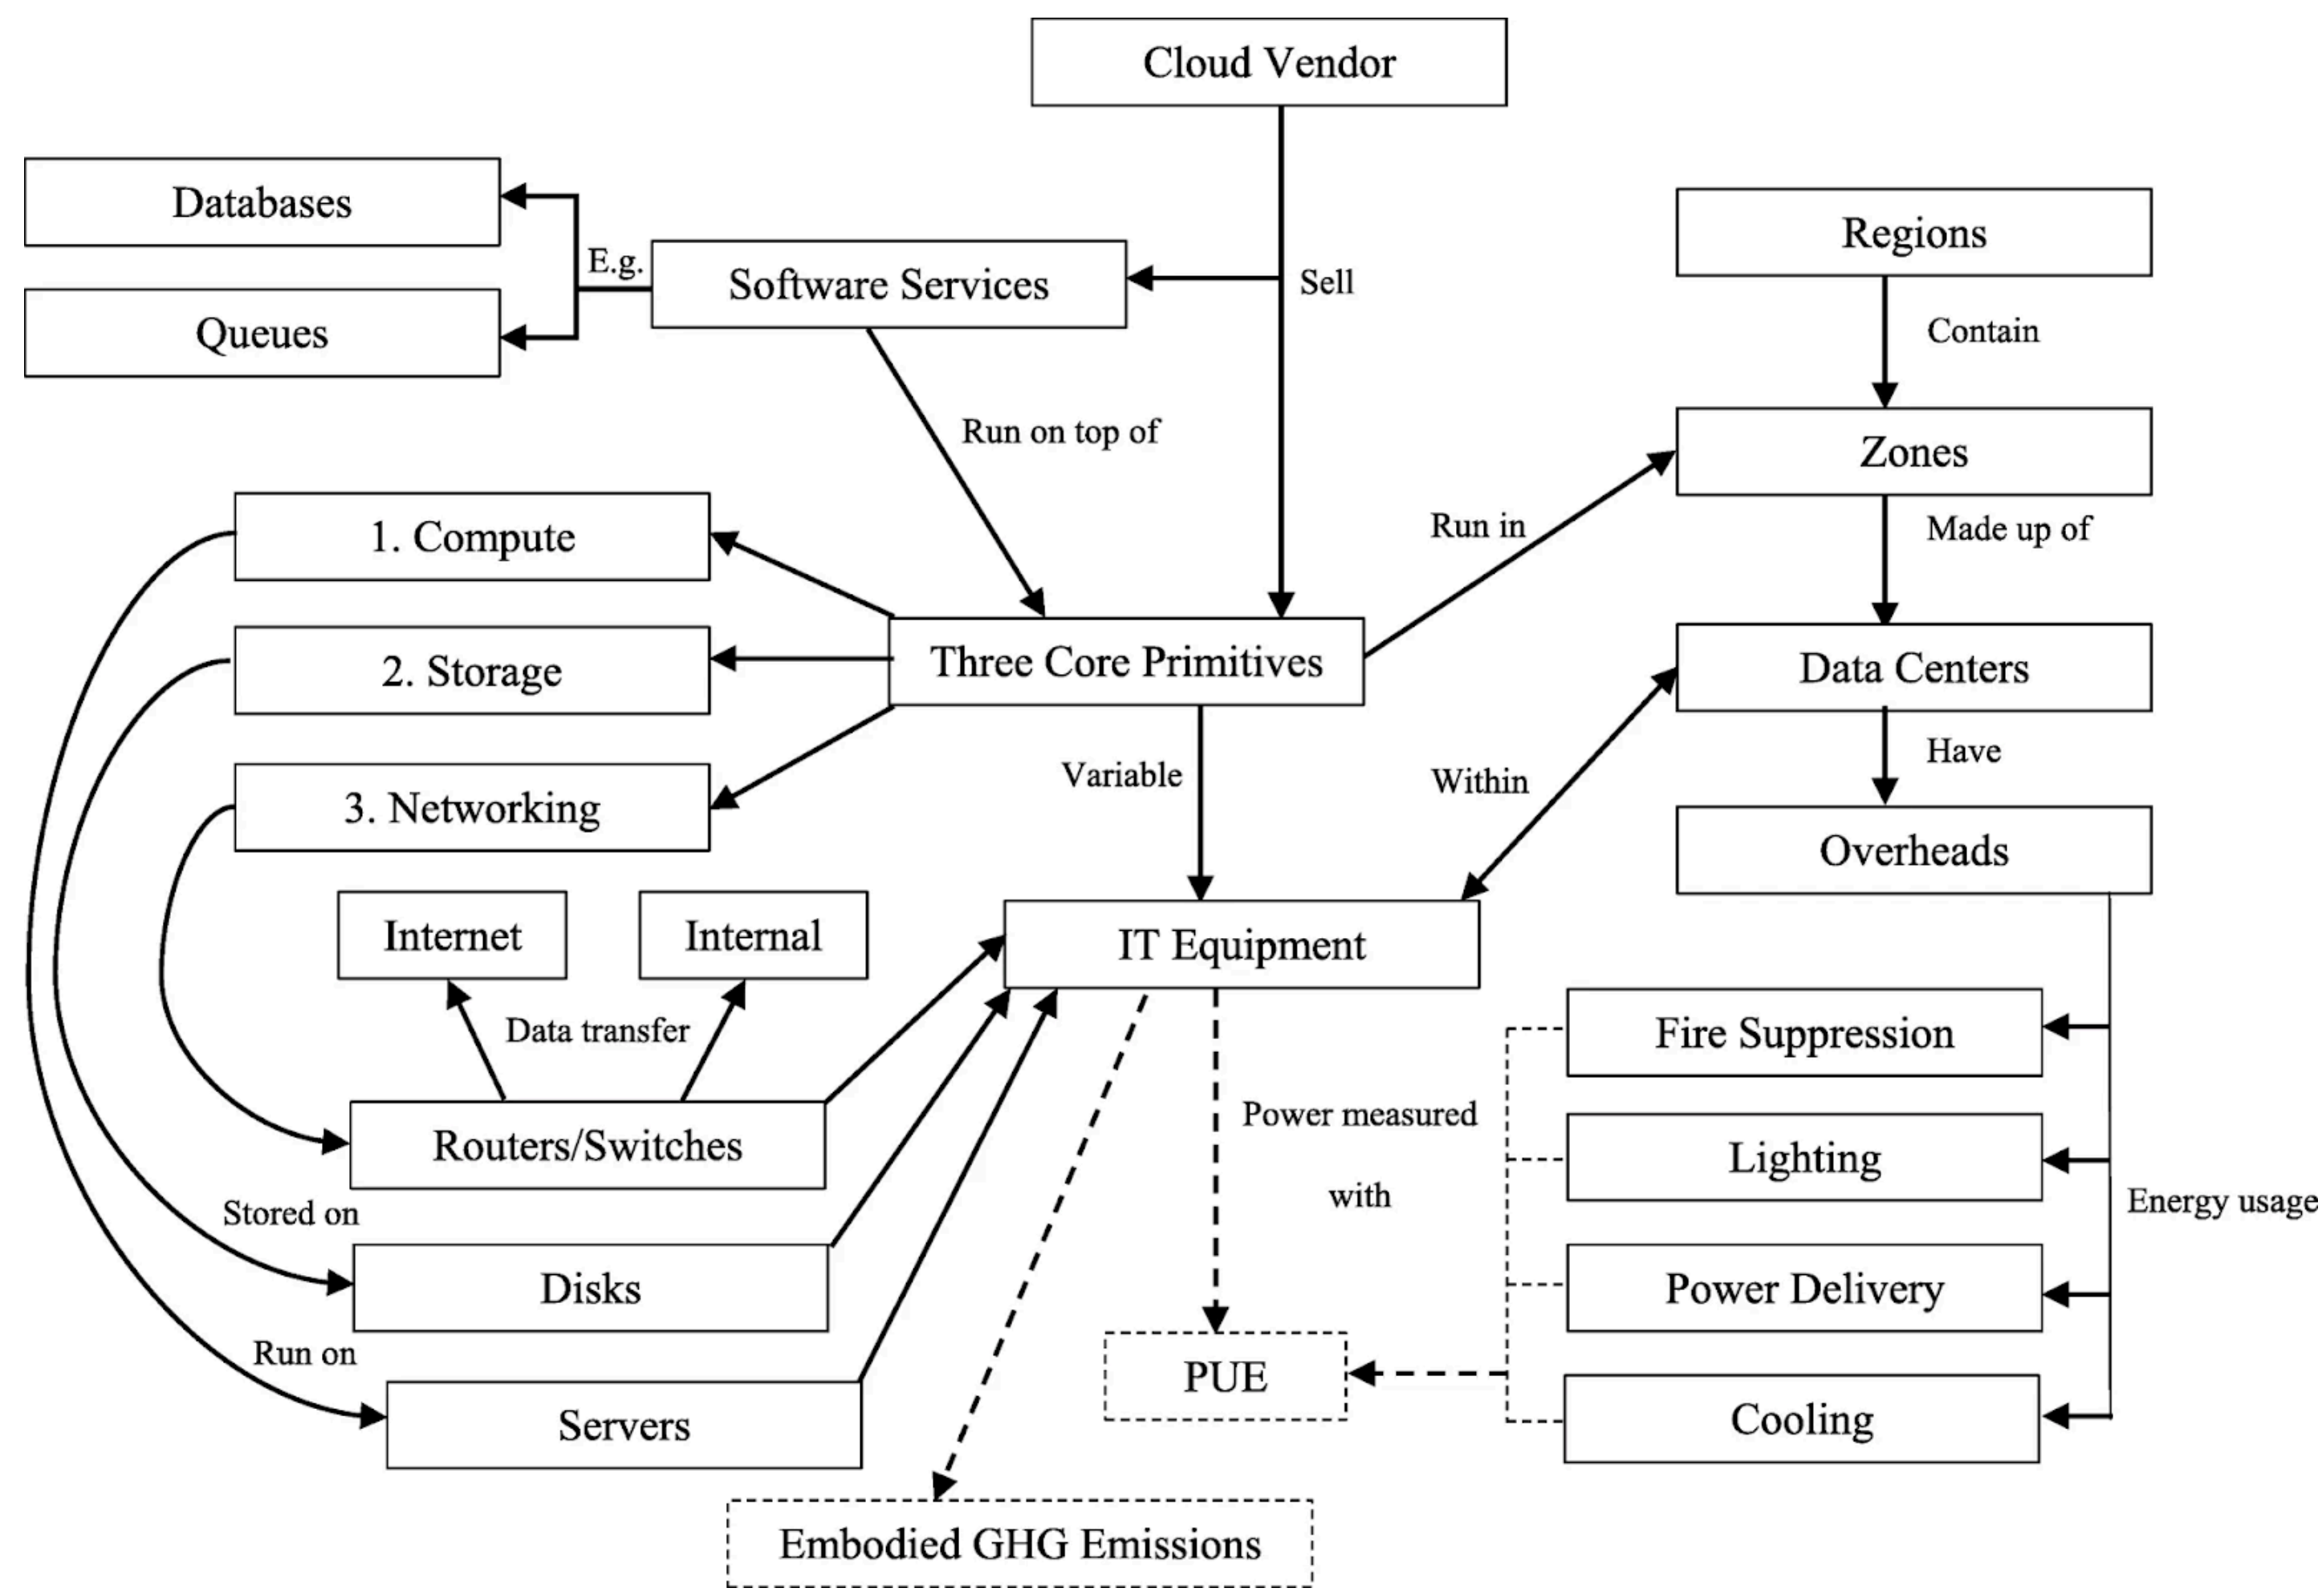
MESURE DE CONSOMMATION ÉLECTRIQUE SUR GPU

- L'interface SMI propose nativement une valeur de consommation lue directement par le driver NVIDIA

```
Every 2.0s: nvidia-smi
Fri Jun 29 11:25:43 2018
+-----+
| NVIDIA-SMI 390.48                  Driver Version: 390.48                |
+-----+-----+
| GPU  Name            Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
+-----+-----+
|  0   Tesla M40             On      | 00000000:03:00.0 Off |             0        |
|  0%   57C    P0     217W / 250W | 10806MiB / 11448MiB |    99%    Default   |
+-----+-----+
|  1   Tesla M40             On      | 00000000:04:00.0 Off |             0        |
|  0%   58C    P0     200W / 250W |  9992MiB / 11448MiB |    99%    Default   |
+-----+-----+
|  2   Tesla M40             On      | 00000000:84:00.0 Off |             0        |
|  0%   56C    P0     222W / 250W |  9955MiB / 11448MiB |    99%    Default   |
+-----+-----+
|  3   Tesla M40             On      | 00000000:85:00.0 Off |             0        |
|  0%   57C    P0     203W / 250W |  9960MiB / 11448MiB |    98%    Default   |
+-----+-----+

+-----+-----+
| Processes:                                     GPU Memory |
|  GPU       PID    Type    Process name      Usage    |
+-----+-----+
|    0      35346     C   python2           110MiB   |
|    0      55323     C    python          10673MiB  |
|    1      55323     C    python           9970MiB   |
+-----+-----+
```

MESURER LA CONSOMMATION ÉLECTRIQUE CHEZ UN CSP



Mesure d'efficacité énergétique

$$PUE = \frac{\text{Total Facility Energy}}{\text{IT Equipment Energy}} = 1 + \frac{\text{Non IT Facility Energy}}{\text{IT Equipment Energy}}$$

Mesure d'efficacité de refroidissement

$$WUE = \frac{\text{Annual Water Usage}}{\text{IT Equipment Energy}}$$

MESURE DANS UN ENVIRONNEMENT VIRTUALISÉ GRÂCE A SCAPHANDRE


End-users



VMs



Hyperviseur




- Données exportées grâce à un agent
- Volume partagé entre toutes les VMS
- Fichier des données lues depuis le RAPL exportées grâce à un Prometheus branché sur l'hyperviseur

CODE CARBON

- ▶ Outil à destination des data scientists & développeurs pythons
- ▶ Lit le fichier RAPL ou ceux émis par l'[Intel Power Gadget](#)
- ▶ Fonctionne sur GPU grâce au driver Nvidia-SMI
- ▶ Données exportées en CSV ou sur une API REST
- ▶ Monitoring de machine / serveur vers l'API

```
from codecarbon import EmissionsTracker

tracker = EmissionsTracker()
tracker.start()
# GPU Intensive code goes here
tracker.stop()
```



Exemplary Equivalents



74.73 %
of weekly
American
household
emissions



293 miles
driven



52 days
of 32-inch
LCD TV
watched

MLCO2 EMISSIONS CALCULATOR

- ▶ Choix du CSP & des data centers
- ▶ Comparaison du hardware
- ▶ Partage des émissions estimées

The screenshot shows the 'ML CO2 Impact' website interface. At the top, there's a navigation bar with links: 'Compute' (highlighted), 'Learn', 'Act', 'About', and 'Authors'. The main heading is 'Machine Learning Emissions Calculator'. Below it, a subheading reads: 'Choose your hardware, runtime and cloud provider to estimate the carbon impact of your research.' A paragraph explains: 'This calculator will give you 2 numbers: the raw carbon emissions produced and the approximate offset carbon emissions. The latter number depends on the grid used by the cloud provider and we are open to update our estimates if anything looks inaccurate or outdated.'

The input section contains four fields: 'Hardware type' (set to 'Titan V'), 'Hours Used' (set to '48'), 'Cloud Provider' (set to 'Google Cloud Platform'), and 'Region of Compute' (set to 'asia-east1'). A red 'COMPUTE' button is centered below these fields.

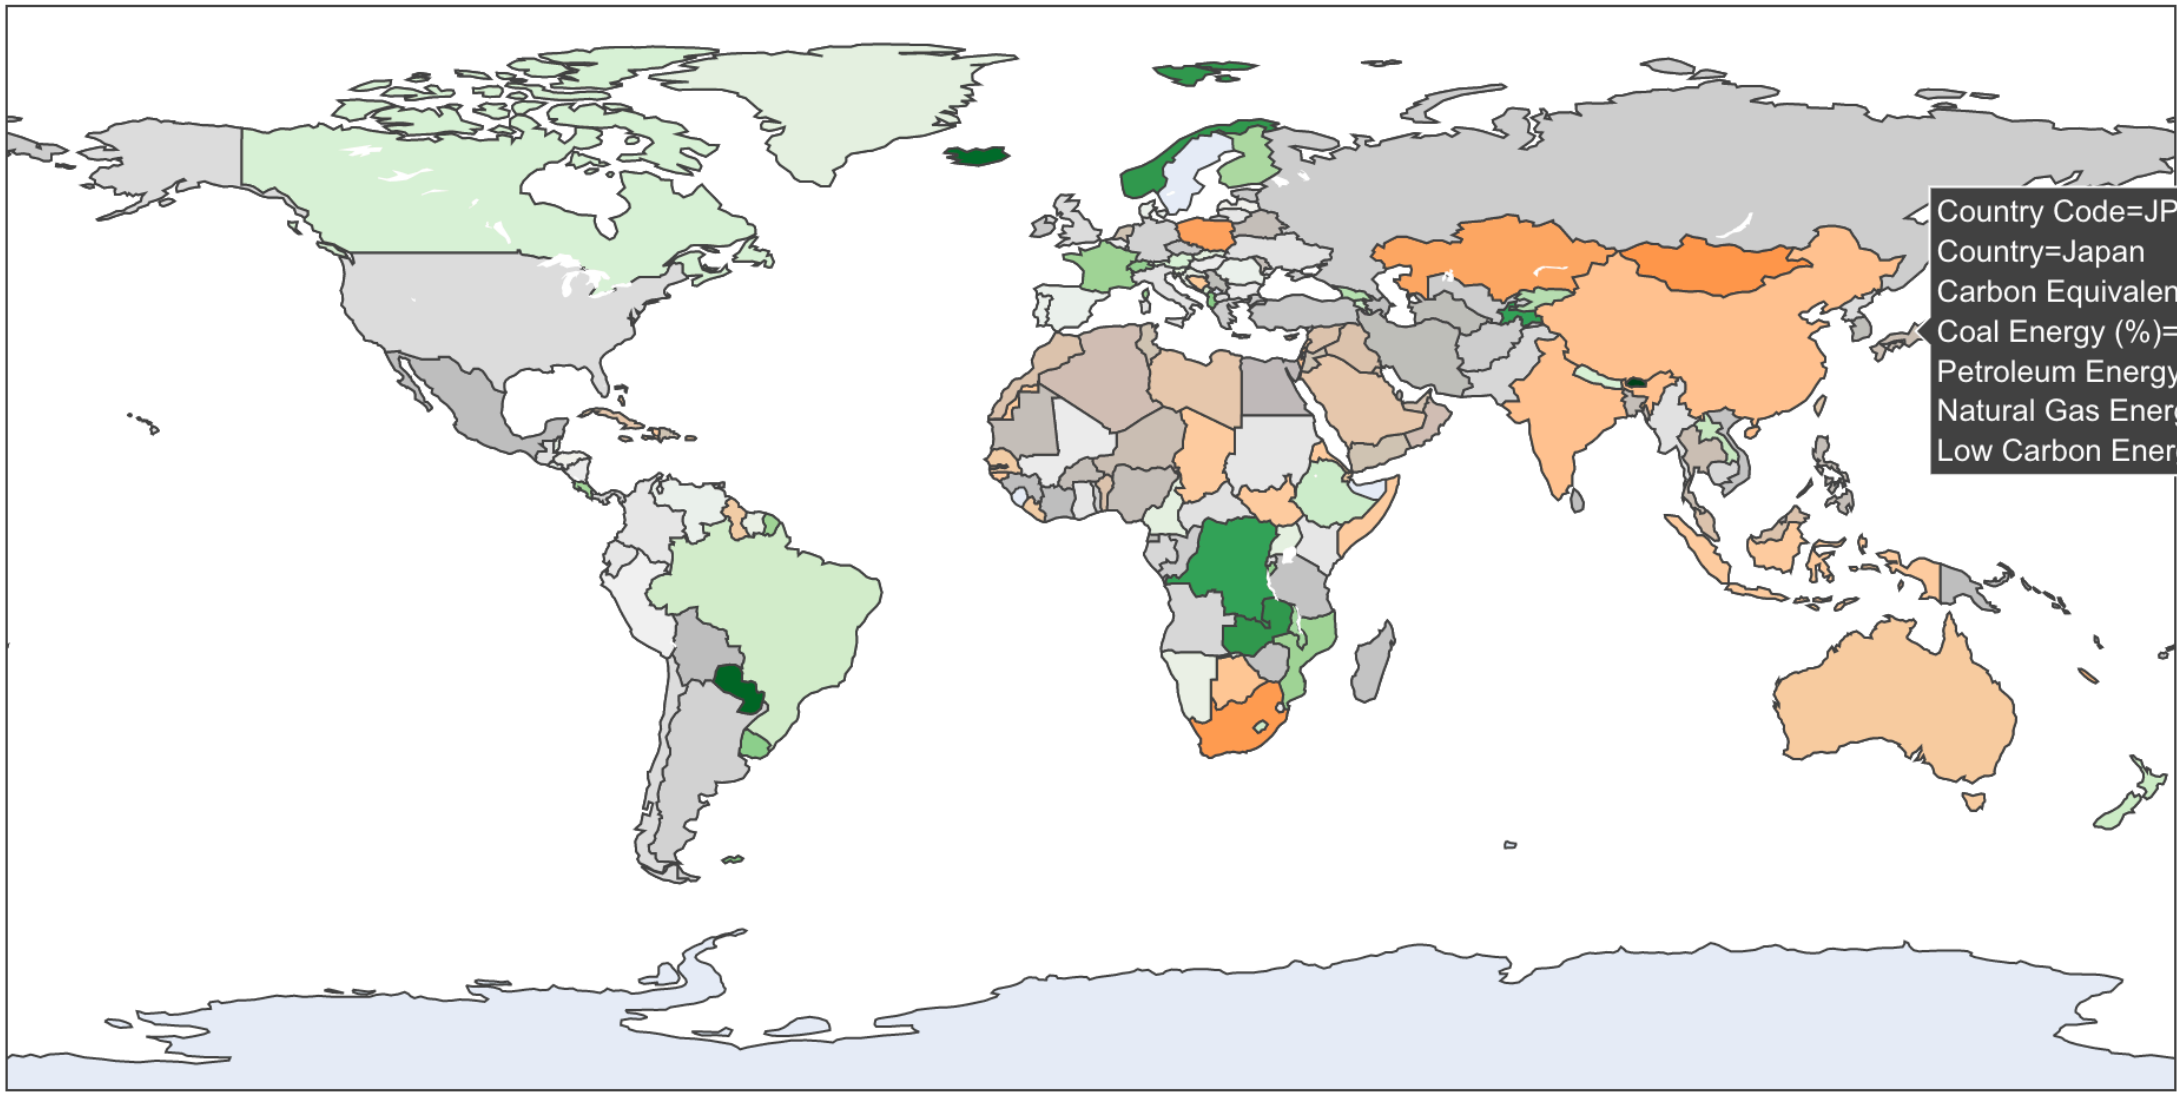
The results are displayed in a white box with a red border. It shows 'CARBON EMITTED' as 6.72 kg CO₂ eq. and 'CARBON ALREADY OFFSET BY PROVIDER' as 6.72 kg CO₂ eq. Below this, it explains the calculation: 'Power consumption x Time x Carbon Produced Based on the Local Power Grid: 250W x 48h = 12 kWh x 0.56 kg eq. CO₂/kWh = 6.72 kg eq. CO₂'. It also provides a comparison: 'Had this model been run in Google Cloud Platform's europe-west6 region, the carbon emitted would have been of 0.19 kg eq. CO₂'. A small 'WHAT CAN I DO?' button is visible on the right side of the results box.

CODE CARBON : ESTIMATION DE L'EQUIVALENT CARBON

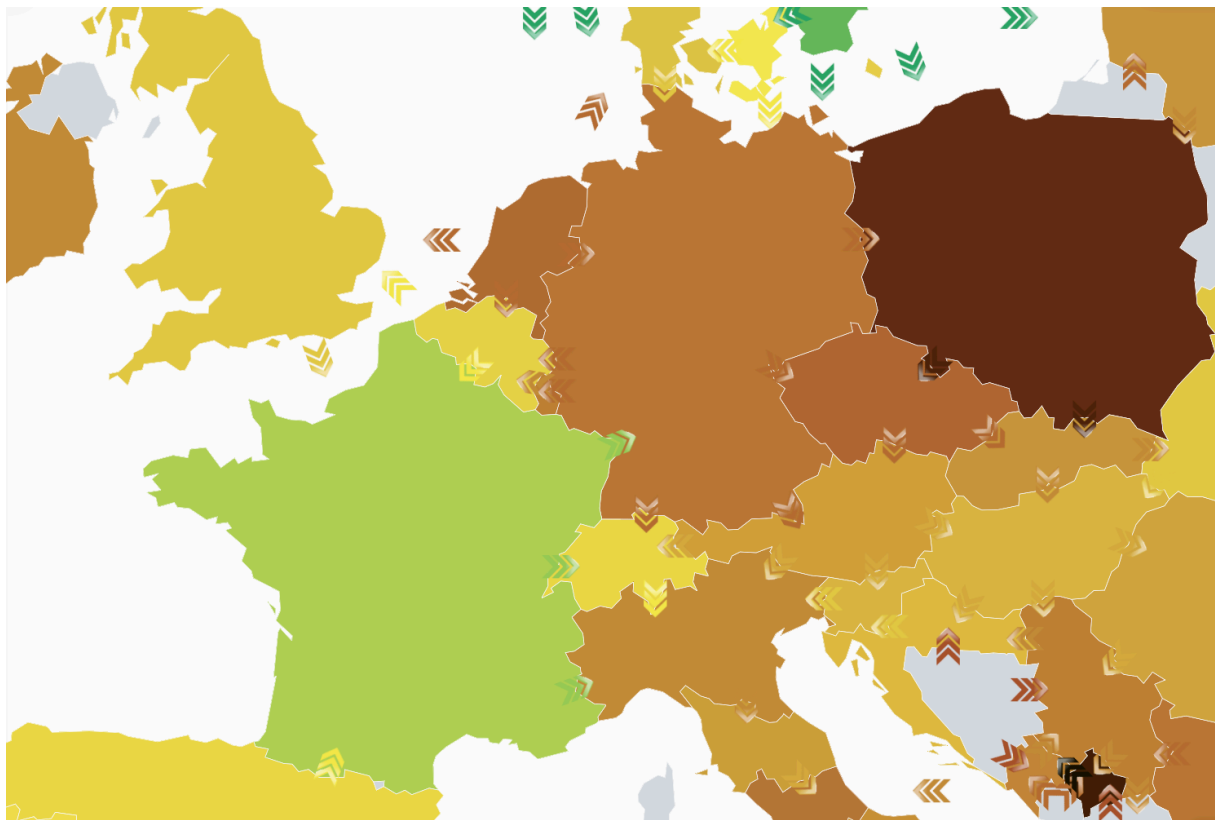

Global Benchmarks

Emissions Equivalent

Energy Mix



Country Code=JPN
Country=Japan
Carbon Equivalent (kg)=86.26749
Coal Energy (%)=24.4
Petroleum Energy (%)=41.3
Natural Gas Energy (%)=24.6
Low Carbon Energy (%)=9.7



Basé sur l'API d'[ElectricityMap](#)

Croise les données de consommation avec celles de mix énergétique de la localisation du traitement

Cartographie les données de PUE des gros cloud providers US

CODE CARBON USERS

divyshah/text-categorization

like

1

Text Classification

PyTorch

Transformers

divyshah/autonlp-data-text-ca

Model card

Files and versions

Model Trained Using AutoNLP

- Problem type: Multi-class Classification
- Model ID: 38039618
- CO2 Emissions (in grams): 214.76275947482927

Status	Name	Tags	Server end time	File name	Duration	user_name	steps
✓	f2466290d		11/11/20 06:43 PM	Jupyter interactive	00:00:48	dsblank	1875

Charts

Panels

Code

Hyper Parameters

Metrics

Graph Definition

Output

System Metrics

Installed Packages

Notes

Graphics

Audio

CodeCarbon Footprint

CODE CARBON

Select an Experiment

f2466290de4e462b9cc9b32476e925

orAnalyze Project

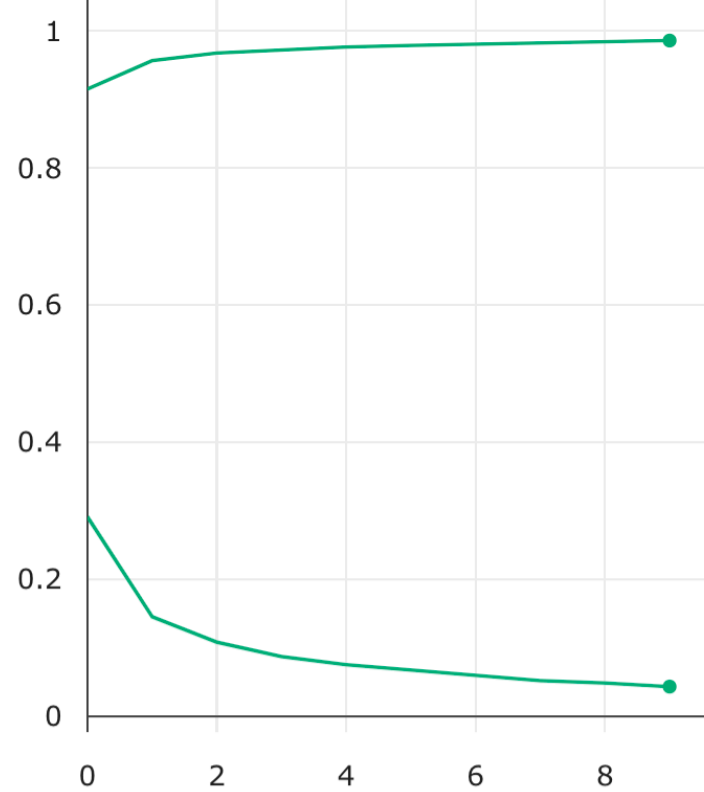
Showing results for experiment:
Infrastructure Hosted at **United States**
Power Consumption : **0.0003 kWh**
Carbon Equivalent : **0.0001 kg**

Exemplary Equivalents

accuracy,loss VS epoch

accuracy

loss





ROADMAP

- ▶ Sortie d'une V2 :
 - ▶ Amélioration de l'estimation (RAM, Windows, méthodes de fallback)
 - ▶ Intégration d'une API & dashboard communautaire
 - ▶ Support de nouvelles puces
- ▶ A venir :
 - ▶ Application web sécurisée et déployable
 - ▶ Amélioration du support des environnements de calcul intensif (HPC, distribués, etc.)



RESOURCES

- ▶ Package [documentation](#) & [repository](#)
- ▶ [Quantifying emissions in Machine Learning](#), V. Schmidt et al., 2019, MILA
- ▶ MLCo2 emissions [calculator](#) & [manifest](#), MILA
- ▶ [Code Carbon API](#)